



Universidad Nacional de San Agustín

Sistema Gestor de Base de Datos

Alberto Gabriel Torres Ara

BASE DE DATOS 2



Estructura del disco

```

v p0
  v f0
    v t0
      s0
      s1
      s2
      s3
      s4
    > t1
    > t2
  > f1
    > t0
    > t1
    > t2
  > p1
  v p2
    v f0
      > t0
      v t1
        s0
        s1
        s2
        s3
        s4
      > t2
    > f1
  > p3
```

```
all23tor@agtavictus:~/projects/UNSA/BBDD2/disco$ ./build/disco
```

El disco aún no existe, se procederá a su creación

Cantidad de platos: 4

Cantidad de pistas por superficie: 3

Cantidad de sectores por pista: 5

Tamaño de cada sector en bytes: 100000

Número de sectores por bloque: 8

"/home/all23tor/projects/UNSA/BBDD2/disco/disk"

└─ "/home/all23tor/projects/UNSA/BBDD2/disco/disk/p0"

└─ "/home/all23tor/projects/UNSA/BBDD2/disco/disk/p0/f0"

└─ "/home/all23tor/projects/UNSA/BBDD2/disco/disk/p0/f0/t0"

└─ "/home/all23tor/projects/UNSA/BBDD2/disco/disk/p0/f0/t0/s0"

└─ "/home/all23tor/projects/UNSA/BBDD2/disco/disk/p0/f0/t0/s1"

└─ "/home/all23tor/projects/UNSA/BBDD2/disco/disk/p0/f0/t0/s2"

└─ "/home/all23tor/projects/UNSA/BBDD2/disco/disk/p0/f0/t0/s3"

└─ "/home/all23tor/projects/UNSA/BBDD2/disco/disk/p0/f0/t0/s4"

└─ "/home/all23tor/projects/UNSA/BBDD2/disco/disk/p0/f0/t1"

Si el disco ya existe,
leer la información
del disco del primer
sector

```
inline void read_disk_info() {  
    auto disk = fs::current_path() / "disk";  
    auto first_path = disk / "p0" / "f0" / "t0" / "s0";  
    ip::file_mapping first_file(first_path.c_str(), ip::mode_t::read_only);  
    ip::mapped_region first_map(first_file, ip::mode_t::read_only);  
    auto data = reinterpret_cast<char*>(first_map.get_address());  
    globalDiskInfo = reinterpret_cast<DiskInfo*>(*data);  
}
```

```
static inline struct DiskInfo {  
    int plates;  
    int tracks;  
    int sectors;  
    int bytes;  
    int block_size;  
} globalDiskInfo = {-1, -1, -1, -1, -1};
```

Si el disco aún no
existe, crearlo y
escribir la
información en el
primer sector

```
auto first_path = disk_path / "p0" / "f0" / "t0" / "s0";  
ip::file_mapping first_file(first_path.c_str(), ip::mode_t::read_write);  
ip::mapped_region first_map(first_file, ip::mode_t::read_write);  
auto data = reinterpret_cast<char*>(first_map.get_address());  
globalDiskInfo = diskInfo;  
reinterpret_cast<DiskInfo*>(*data) = globalDiskInfo;  
  
fs::create_directory(blocks_path);  
}
```

Almacenamiento

Registros de longitud fija

Esquema

```
90 Address* write_header(const std::string& table_name,  
91                       const std::vector<Db::Column>& columns) {  
92     auto first_data = CowBlock::load_writeable_sector({0}) + sizeof(DiskInfo);  
93     auto tables = reinterpret_cast<Table*>(first_data);  
94     int table_idx = 0;  
95  
96     while (tables[table_idx].name[0] != '\0')  
97       | table_idx++;  
98  
99     auto& table = tables[table_idx];  
100    for (int i = 0; i < table_name.size(); i++)  
101      | table.name[i] = table_name[i];  
102    for (int i = table_name.size(); i < table.name.size(); i++)  
103      | table.name[i] = '\0';  
104  
105    auto header_sector = request_empty_sector();  
106    auto header_data = CowBlock::load_writeable_sector(header_sector);  
107    table.sector = header_sector;  
108  
109    auto records_start = reinterpret_cast<Address*>(header_data);  
110    *records_start = NullAddress;  
111    header_data += sizeof(Address);  
112
```



Almacenamiento

Registros de longitud fija

Información en el esquema

Tabla	Registro inicial	Columnas
	Columna 0	Tipo 0
	Columna 1	Tipo 1
	Columna 2	Tipo 2
	Columna 3	Tipo 3
	Columna 4	Tipo 4
	Columna 5	Tipo 5


```
105 auto header_sector = request_empty_sector();
106 auto header_data = CowBlock::load_writeable_sector(header_sector);
107 table.sector = header_sector;
108
109 auto records_start = reinterpret_cast<Address*>(header_data);
110 *records_start = NullAddress;
111 header_data += sizeof(Address);
112
113 int column_size = columns.size();
114 for (char c : pun_cast(column_size))
115     |(header_data++) = c;
116
117 for (auto& column : columns)
118     |for (char c : pun_cast(column))
119         |(header_data++) = c;
120
121 return records_start;
122 }
```


Almacenamiento

Registros de longitud fija

Sector en uso

```
int recods_per_sector =
    (globalDiskInfo.bytes - sizeof(Address) - sizeof(int)) / record_size;
*records_start = request_empty_sector();
auto record_data = CowBlock::load_writeable_sector(*records_start);
auto next_sector = reinterpret_cast<Address*>(record_data);
record_data += sizeof(Address);
auto record_count = reinterpret_cast<int*>(record_data);
record_data += sizeof(int);
*next_sector = NullAddress;
*record_count = 0;

while (std::getline(file, line)) {
    if (*record_count == recods_per_sector) {
        *next_sector = request_empty_sector();
        record_data = CowBlock::load_writeable_sector(*next_sector);
        next_sector = reinterpret_cast<Address*>(record_data);
        record_data += sizeof(Address);
        record_count = reinterpret_cast<int*>(record_data);
        record_data += sizeof(int);
        *next_sector = NullAddress;
        *record_count = 0;
    }
}
```

Sgte sector

Cantidad de registros

Registro 0

Registro 1

Registro 2


Registro 3

Registro 4

Registro 5

...

Bloques



Bloques en
vertical

Minimiza el
movimiento del
brazo

Facilita la lógica
de solicitar
sectores vacíos



Direccinamiento de bloques



Bloque 0 Bloque 1 Bloque 2 Bloque 3

Plato 0

Plato 1

Plato 2

Plato 3

	Sector 0		Sector 4		Sector 8		Sector 12
	Sector 1		Sector 5		Sector 9		Sector 13
	Sector 2		Sector 6		Sector 10		Sector 14
	Sector 3		Sector 7		Sector 11		Sector 15



De lugar físico a dirección lógica y viceversa

```
struct Address {
    int address;

    bool operator==(const Address&) const = default;
    static Address from_tree(int plate, int surface, int track, int sector) {
        return {plate +
                globalDiskInfo.plates *
                (sector + globalDiskInfo.sectors *
                 (track + globalDiskInfo.tracks * surface))};
    }
}
```

```
fs::path to_path() const {
    int address = this->address;
    auto plate = address % globalDiskInfo.plates;
    address /= globalDiskInfo.plates;
    auto sector = address % globalDiskInfo.sectors;
    address /= globalDiskInfo.sectors;
    auto track = address % globalDiskInfo.tracks;
    address /= globalDiskInfo.tracks;
    auto surface = address % 2;

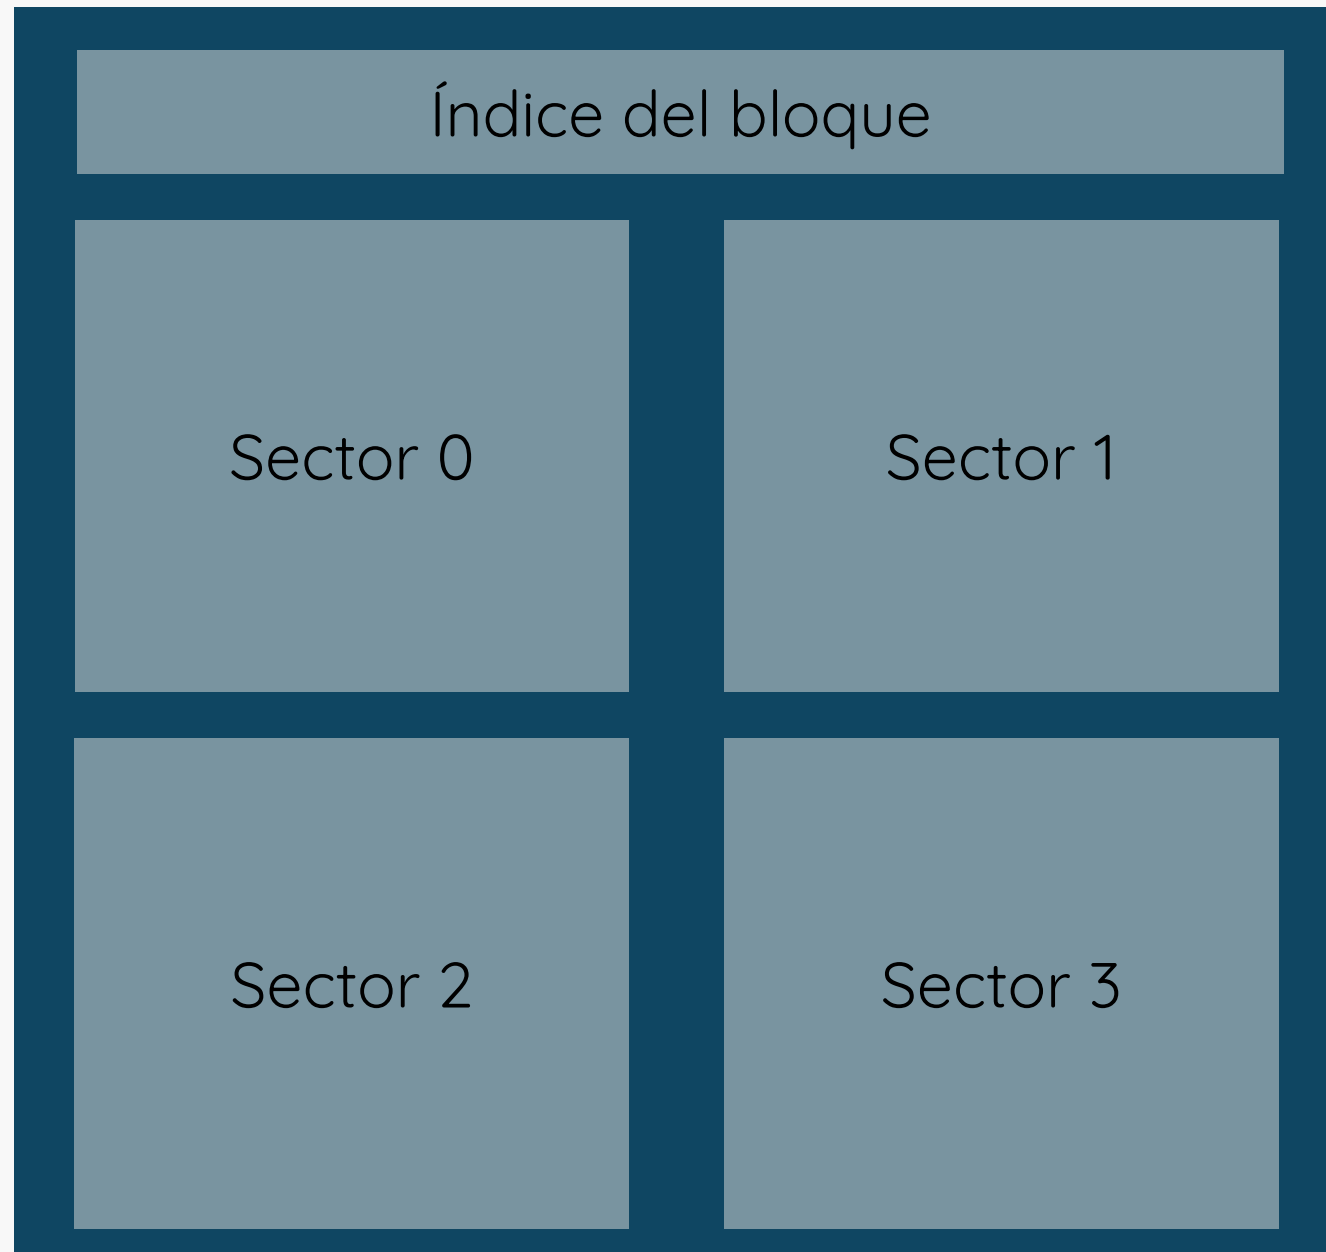
    return disk_path / ('p' + std::to_string(plate)) /
        ('f' + std::to_string(surface)) / ('t' + std::to_string(track)) /
        ('s' + std::to_string(sector));
}
```

Bloque Copy on Write

```
class CowBlock {
    mutable int written = false;
    ip::mapped_region block_map;

    static inline std::set<CowBlock, std::less<>> loaded_blocks;
```

Estructura



De índice de bloque a direcciones de sectores

```
int block_idx = *reinterpret_cast<const int*>((data() - sizeof(int)));
for (int sector = 0; sector < globalDiskInfo.block_size; sector++) {
    auto sector_data = data() + sector * globalDiskInfo.bytes;
    Address sector_address = {block_idx * globalDiskInfo.block_size + sector};
```

Solo sobrescribimos los sectores si escribimos en el bloque

```
~CowBlock() {
    if (!written)
        return;
```

Guardamos los bloques cargados para evitar volver a cargar el mismo sector

[illegible]

Referencias

1. Silberschatz, A., Korth, H. F., & Sudarshan, S. (2011). Database system concepts (6th ed.).
2. McGraw-Hill. Garcia-Molina, H., Ullman, J. D., & Widom, J. (2008). Database systems: The complete book (2nd ed.). Pearson Education.
3. Gamma, E., Helm, R., Johnson, R., & Vlissides, J. (1995). Patrones de diseño: Elementos de software orientado a objetos reutilizable (1.ª ed.). Addison-Wesley Iberoamericana.

.....

*Muchas
gracias*

.....